

B⁺ Trees

9/10

What makes a good index?

- useful
- small
- make finding quicker

2 types of indices:

- check existence: hash
- range queries: ordered index

Things to consider

- Access type: single elt vs range
- Access time: how long to find something
- Insertion time:
- Deletion time:
- Space

def **search key** as attribute used to lookup a record


About B⁺ tree:

- balanced tree w/ every path from the root to a leaf has same length
- each non-leaf node have between $\lceil \frac{n}{2} \rceil$ and n children (where n is fixed for the tree)

Structure of a leaf

- up to $n-1$ search keys
 K_1, \dots, K_{n-1}
w/ $K_i < K_j$ if $i < j$
- up to n pointers P_1, \dots, P_n
w/ P_i points to a record w/ search key K_i
- each leaf points to "neighbor" leaf

$n=4$



Records:

id	name	dept	Course #
20	Sh	CS	447
21	Mu	CS	451
21	Fa	CS	432
23	Pa	CS	127
24	Mu	CS	540


- each leaf may contain
between $\lceil \frac{n-1}{2} \rceil$ and $n-1$ values

Non leaf nodes

- similar struct to leaves

But:

- ptrs are to other nodes
- no ptr to sibling
- each node holds between $\lceil \frac{n}{2} \rceil$ and n ptrs
- # of ptrs sometimes called **fanout**



Bin search tree vs B-tree

consider table w/ 1,000,000 records

Bin search tree:	height	$\log_2(1,000,000) \approx 20$
B-tree $n=100$:	height	$\log_{50}(1,000,000) \approx 4$