Sorting (external)

9/17

External merge sort

Let M be the # of blocks in main memory buffer for sorting

① Fill buffer, sort buffer, write buffer
each iteration is called a run

```
i = 0;
repeat
    read M blocks of the relation, or the rest of the relation,
        whichever is smaller;
    sort the in-memory part of the relation;
    write the sorted data to run file R_i;
    i = i + 1;
until the end of the relation
```

② merge step

Let N be the # of runs that we perform

assure $N < M$

$L = 3$

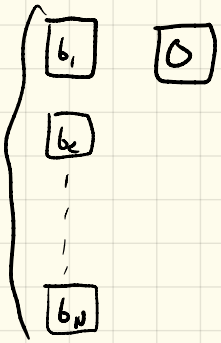

the hard case
$N >> M$

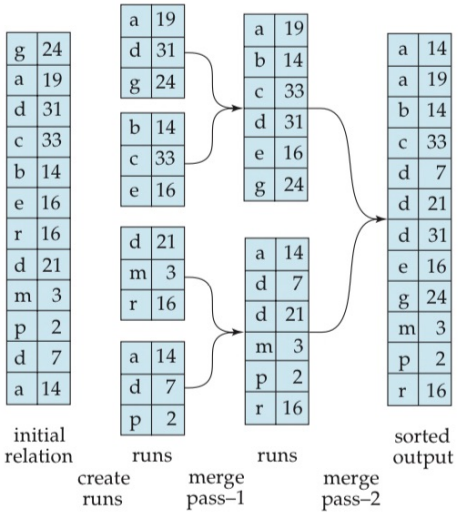

Figure 12.4 External sorting using sort–merge.

each pass reduces $M-1$ runs
← we need 1-block for output buffer

read one block of each of the N files R_i into a buffer block in memory;
repeat
    choose the first tuple (in sort order) among all buffer blocks;
    write the tuple to the output, and delete it from the buffer block;
    if the buffer block of any run R_i is empty **and not** end-of-file(R_i)
        then read the next block of R_i into the buffer block;
until all input buffer blocks are empty

Analysis of external merge sort
(disk access)

b_r is # of blocks containing relation r

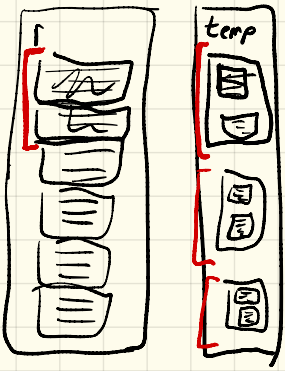

① read all blocks
(sort)
write all blocks
$\rightarrow 2b_r$ transfers

② # of runs $\lceil b_r/m \rceil$

③ # of merge passes $\lceil \log_{m-1}(b_r/m) \rceil$

④ each merge pass reads and writes all blocks
($2b_r$ blocks)

$\Rightarrow b_r\left(2\lceil \log_{m-1}(b_r/m)\rceil + 1\right)$ ≤ not counting the final write

* Account for # of seeks