

Column Store

9/27

HBase as an example:

- non-relational
- made for big-jobs (terabytes of data and 10s of nodes)

HBase history

- Apache project
based on

"BigTable": A distributed Storage System for Structured Data"

Unique Features:

- scalable
- distributed config (no central point of failure)
- tolerate node failure in a large cluster
- versions data
- compression
- garbage collection (exp. data)
- consistency and partition tolerance (CP)

Data model

- maps of maps

key (string) → "row" of data

map w/ key called **columns**

grouped into **column families**
values are byte arrays

	row keys	column family "color"	column family "shape"
row	"first"	"red": "#F00" "blue": "#00F" "yellow": "#FF0"	"square": "4"
row	"second"		"triangle": "3" "square": "4"